

COMP755-Lect03

September 3, 2018

1 COMP 755

Plan for today

1. Review likelihood, log-likelihood, maximum likelihood estimates
2. Compute Maximum Likelihood Estimate for a Gaussian model
3. Introduce Linear Regression
4. Introduce Gradient Ascent/Descent methods

2 Last time -- distributions

Last time we reviewed probability distributions of one dimensional random variables: * Bernoulli (coin toss), Binomial (count of heads in multiple coin tosses) * Categorical (die roll), Multinomial (count of different die sides in multiple rolls)

Continuous distributions: * Gaussian * Laplace

3 Last time -- likelihood

We introduced **likelihood** function.

$$\mathcal{L}(\theta|\mathbf{x}) = p(\underbrace{\mathbf{x}}_{\text{Data}} | \underbrace{\theta}_{\text{parameters}}) = \prod_i p(x_i|\theta)$$

It tells us how well our model fits the data for particular parameters.

An example of model

$$p(x|\theta) = \theta^x(1-\theta)^{1-x}$$

data

$$\mathbf{x} = \{0, 1, 1, 1\}$$

and likelihood

$$\mathcal{L}(\theta|\mathbf{x}) = \prod_i p(x_i|\theta) = p(x_1|\theta)p(x_2|\theta)p(x_3|\theta)p(x_4|\theta) = (1-\theta)\theta^3$$

4 Last time -- log-likelihood and MLE

We introduced log-likelihood function

$$\log \mathcal{L}(\theta | \text{Data}).$$

An example: Given data $\text{Data} = \{0, 0, 0, 1\}$ and model $p(x|\theta) = \theta^x(1-\theta)^{1-x}$

$$-2.25 = \log \mathcal{L}(0.25 | \mathbf{x}) > \log \mathcal{L}(0.99 | \text{Data}) = -13.83$$

```
In [1]: import numpy
        from __future__ import print_function
        def likelihood(theta,xs):
            p = 1.0
            for x in xs:
                p = p*theta**x * (1-theta)**(1-x)
            return p

        xs = [0,0,0,1]
        theta1 = 0.99
        theta2 = 0.25
        loglik1 = numpy.log(likelihood(theta1,xs))
        loglik2 = numpy.log(likelihood(theta2,xs))
        print("Log-Likelihood(",theta1,"|",xs, ")=",loglik1)
        print("Log-Likelihood(",theta2,"|",xs, ")=",loglik2)
```

Log-Likelihood(0.99 | [0, 0, 0, 1])= -13.8255608938

Log-Likelihood(0.25 | [0, 0, 0, 1])= -2.24934057848

5 Last time -- finding Maximum Likelihood Estimate (MLE)

We talked about finding a maximizer of the log-likelihood called **Maximum Likelihood Estimate (MLE)**

$$\theta^{\text{MLE}} = \arg \max_{\theta} \log \mathcal{L}(\theta | \text{Data})$$

which can be interpreted as the parameter for which the data is most probable under the model.

We recalled that in order to maximize a function, $f(x)$, we find x for which

$$\frac{\partial}{\partial x} f(x) = 0.$$

We used this approach to find MLE for a Bernoulli model

$$p(x|\theta) = \theta^x(1-\theta)^{1-x}$$

and came up with

$$\theta^{\text{MLE}} = \frac{\overbrace{\sum_i [x_i = 1]}^{\text{count of 1s in data}}}{\underbrace{\sum_i [x_i = 1]}_{\text{count of 1s in data}} + \underbrace{\sum_i [x_i = 0]}_{\text{count of 0s in data}}}$$

6 Finding μ^{MLE} of Gaussian distribution

We left as an exercise a problem to come up with a maximum likelihood estimate for parameter μ of a Gaussian distribution

$$p(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$$

So we will do that now.

Likelihood function is

$$\mathcal{L}(\mu, \sigma^2|\mathbf{x}) = \prod_{i=1}^N p(x_i|\mu, \sigma^2) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x_i-\mu)^2}$$

Log-likelihood function is

$$\log \mathcal{L}(\mu, \sigma^2|\mathbf{x}) = \log \prod_{i=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x_i-\mu)^2} = \sum_{i=1}^N \left[-\frac{1}{2} \log 2\pi\sigma^2 - \frac{1}{2\sigma^2}(x_i - \mu)^2 \right]$$

7 Finding μ^{MLE} of Gaussian distribution

Our recipe is: 1. Take the function you want to maximize:

$$f(\mu) = \sum_{i=1}^N \left[-\frac{1}{2} \log 2\pi\sigma^2 - \frac{1}{2\sigma^2}(x_i - \mu)^2 \right]$$

2. Compute its first derivative: $\frac{\partial}{\partial \mu} f(\mu)$ 3. Equate that derivative to zero and solve: $\frac{\partial}{\partial \mu} f(\mu) = 0$

The first derivative is

$$\frac{\partial}{\partial \mu} f(\mu) = \sum_{i=1}^N \left[-\frac{1}{\sigma^2}(x_i - \mu) \right]$$

We equate it to zero and solve

$$\sum_{i=1}^N \left[-\frac{1}{\sigma^2}(x_i - \mu) \right] = 0$$

8 Finding μ^{MLE} of Gaussian distribution

$$\begin{aligned}\sum_{i=1}^N \left[-\frac{1}{\sigma^2} (x_i - \mu) \right] &= 0 \\ \sum_{i=1}^N \frac{1}{\sigma^2} (x_i - \mu) &= 0 \\ \sum_{i=1}^N (x_i - \mu) &= 0 \\ \sum_{i=1}^N x_i &= \sum_i \mu \\ \sum_{i=1}^N x_i &= \mu \sum_i 1 \\ \frac{\sum_{i=1}^N x_i}{\sum_{i=1}^N 1} &= \mu \\ \frac{\sum_{i=1}^N x_i}{N} &= \mu\end{aligned}$$

9 Using MLE estimates in prediction

What can we do with maximum likelihood estimates? We can predict outcome of the next experiment: * Given data $\{1, 0, 1, 1, 0, 0, 0, 0\}$ and assuming Bernoulli model, θ^{MLE} is $\frac{3}{3+5} = \frac{3}{8}$. Hence, probability that the next experiment will yield $x = 1$ is $\frac{3}{8}$

- Given heights of 6th graders in inches $\{56, 57, 59, 63, 55, 61\}$ and assuming Gaussian model, μ^{MLE} is $\frac{56+57+59+63+55+61}{6} = 58.5$ and standard deviation is 3.082. Hence, probability that the next 6th grader will be of height $58.5 \pm 2*3.082$ is 0.95.

Given realizations of the same random variable we can estimate what the next draw might look like.

Q: Suppose you wanted to predict a student's grade in COMP 755? What would you do?

10 Linear Regression

One of the simplest examples of supervised learning is **linear regression**.

An underlying assumption in **linear regression** is that the target variable, y , varies around a mean which is a linear combination of some set of features.

For example, we can try to predict a student's grade in COMP755 class using grades from prereq's and number of beers the student has each week

$$\begin{aligned}\text{COMP755} &= 0.25 + 0.2 * \text{COMP410} + 0.3 * \text{MATH233} + 0.5 * \text{STOR435} - 0.1 * \text{beers} + \epsilon \\ \epsilon &\sim \mathcal{N}(0, 0.1)\end{aligned}$$

Formally, we would write

$$\begin{aligned}y &= \beta_0 + \sum_j x_j \beta_j + \epsilon \\ \epsilon &\sim \mathcal{N}(0, \sigma^2)\end{aligned}$$

or more compactly

$$y|\mathbf{x} \sim \mathcal{N}\left(\beta_0 + \sum_j x_j \beta_j, \sigma^2\right)$$

11 Linear Regression

Taking a closer look at the model:

$$y|\mathbf{x} \sim \mathcal{N}\left(\beta_0 + \sum_{j=1}^p x_j \beta_j, \sigma^2\right)$$

* y is a target variable we are modeling * \mathbf{x} is a vector of p features (aka predictors and covariates)
 * β_0 is a **bias** which does not depend on the features * β_1, \dots, β_p is a vector of weights (one per feature) * σ^2 is variance -- how far can we expect y to be away from $\beta_0 + \sum_{j=1}^p x_j \beta_j$

Q: Which of these items above are parameters? Hint: Which values can we adjust to get a better prediction?

$$\text{COMP755} = 0.25 + 0.2 * \text{COMP410} + 0.3 * \text{MATH233} + 0.5 * \text{STOR435} - 0.1 * \text{beers} + \epsilon$$

12 Linear Regression

Probability of target variable y

$$p(y|\mathbf{x}, \beta_0, \beta, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{1}{2\sigma^2} \left(y_i - \underbrace{\left(\beta_0 + \sum_j x_j \beta_j \right)}_{\text{mean of the Gaussian}} \right)^2 \right\}$$

In the case of the 6th grader's height, we made **the same** prediction for any other 6th grader (58.5 inches).

In our COMP 755 grade example, we compute a potentially different mean for every student.

$$\beta_0 + \beta_{\text{COMP410}} * \text{COMP410} + \beta_{\text{MATH233}} * \text{MATH233} + \beta_{\text{STOR435}} * \text{STOR435} + \beta_{\text{beers}} * \text{beers}$$

13 Linear Regression -- toy example

We can try to fit this model to some data.

Note that these are **not** real UNC student data.

$x_1 = \text{COMP410}$	$x_2 = \text{MATH233}$	$x_3 = \text{STOR345}$	$x_4 = \text{\#beers}$	$y = \text{COMP755}$
3	3	4	0	4
3	3	3	0	3
3	2	2	5	2
2	2	4	0	3
3	3	4	4	3

$x_1 = \text{COMP410}$	$x_2 = \text{MATH233}$	$x_3 = \text{STOR345}$	$x_4 = \text{\#beers}$	$y = \text{COMP755}$
4	4	3	0	4
2	2	4	0	3
2	2	2	0	2
3	4	2	5	3
2	2	3	0	3
2	4	2	0	3
3	3	3	0	3
2	2	4	5	2
4	3	2	6	2
3	4	4	4	3
3	2	2	6	2
3	3	3	0	3
3	2	4	0	3
3	3	4	4	3
3	4	2	0	3

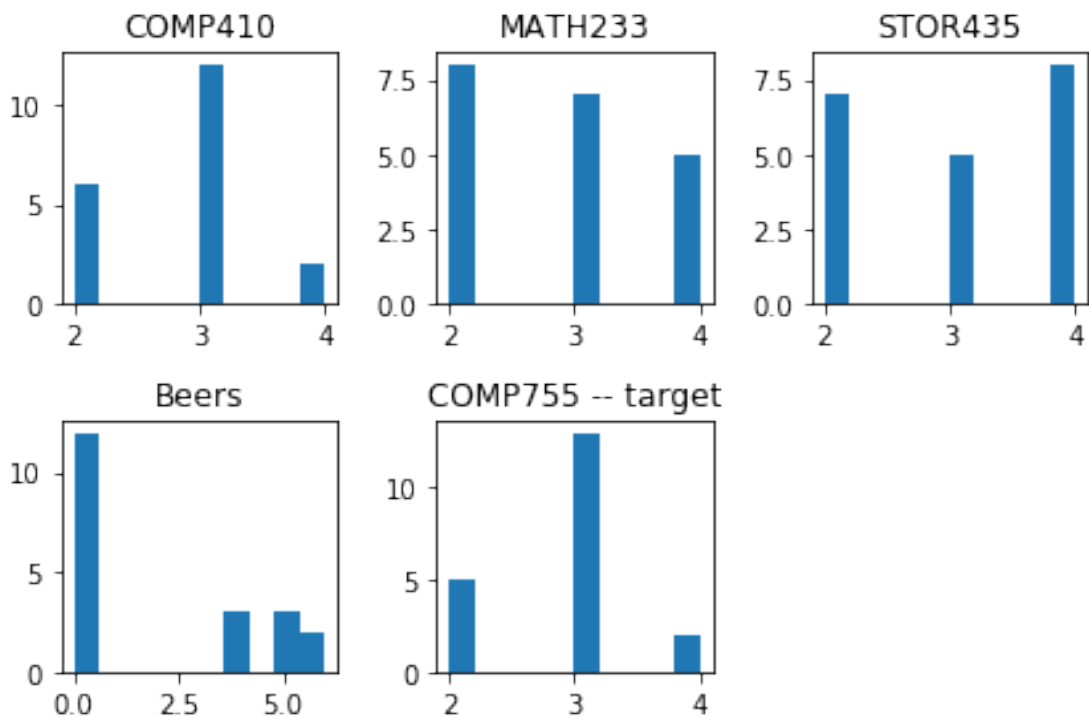
```
In [3]: X = numpy.asarray([[3.0,3.0,4.0,0.0],[3.0,3.0,3.0,0.0],[3.0,2.0,2.0,5.0],
    [2.0,2.0,4.0,0.0],[3.0,3.0,4.0,4.0],[4.0,4.0,3.0,0.0],
    [2.0,2.0,4.0,0.0],[2.0,2.0,2.0,0.0],[3.0,4.0,2.0,5.0],
    [2.0,2.0,3.0,0.0],[2.0,4.0,2.0,0.0],[3.0,3.0,3.0,0.0],
    [2.0,2.0,4.0,5.0],[4.0,3.0,2.0,6.0],[3.0,4.0,4.0,4.0],
    [3.0,2.0,2.0,6.0],[3.0,3.0,3.0,0.0],[3.0,2.0,4.0,0.0],
    [3.0,3.0,4.0,4.0],[3.0,4.0,2.0,0.0]])
```

```
y = numpy.asarray([4.0,3.0,2.0,3.0,3.0,4.0,3.0,2.0,3.0,
    3.0,3.0,3.0,2.0,2.0,3.0,2.0,3.0,3.0,3.0,3.0])
```

```
In [4]: import matplotlib.pyplot as plt
    %matplotlib inline
    import numpy
    features = ['COMP410', 'MATH233', 'STOR435', 'Beers']
    for j in range(X.shape[1]):
        plt.subplot(2,3,j+1)
        plt.hist(X[:,j])
        plt.title(features[j])

    plt.subplot(2,3,5)
    print(y)
    plt.hist(y)
    plt.title('COMP755 -- target')
    plt.tight_layout() # ensure sensible layout of subplots
```

```
[ 4.  3.  2.  3.  3.  4.  3.  2.  3.  3.  3.  3.  2.  2.  3.  2.  3.  3.
  3.  3.]
```



14 Linear Regression -- toy example

Our goal is to fit the model that predicts COMP 755 grade.

Our prediction is based on

$$\text{COMP562} = \beta_0 + \beta_{\text{COMP410}} * \text{COMP410} + \beta_{\text{MATH233}} * \text{MATH233} + \beta_{\text{STOR435}} * \text{STOR435} + \beta_{\text{beers}} * \text{beers} + \epsilon$$

In order to make a prediction we need β_0 -- bias or grade you get regardless of your other grades $\beta_{\text{COMP410}} * \beta_{\text{MATH233}} * \beta_{\text{STOR435}} * \beta_{\text{beers}}$

Note that having σ^2 would enable us to give a range of grades that cover 95% of the probability.

15 Linear Regression -- likelihood

We start by writing out a likelihood for linear regression is

$$\mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) = \prod_{i=1}^N p(y_i | \mathbf{x}_i, \beta_0, \beta, \sigma^2) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{1}{2\sigma^2} \left(y_i - (\beta_0 + \sum_j x_j \beta_j) \right)^2 \right\}$$

Log-likelihood for linear regression is

$$\log \mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) = \sum_{i=1}^N \left[-\frac{1}{2} \log 2\pi\sigma^2 - \frac{1}{2\sigma^2} \left(y_i - (\beta_0 + \sum_j x_j \beta_j) \right)^2 \right].$$

Our goal is still to find β_0, β such that

$$\begin{aligned}\frac{\partial}{\partial \beta_0} \log \mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) &= 0 \\ \frac{\partial}{\partial \beta_1} \log \mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) &= 0 \\ &\dots \\ \frac{\partial}{\partial \beta_p} \log \mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) &= 0\end{aligned}$$

because that will guarantee the parameters can not be further changed to improve the likelihood.

16 Introducing gradient ascent

Previously, we solved equations of type $\frac{\partial}{\partial \theta} \log \mathcal{L}(\theta | \text{Data}) = 0$ in a closed-form.

Here, we will develop a different approach using numerical optimization.

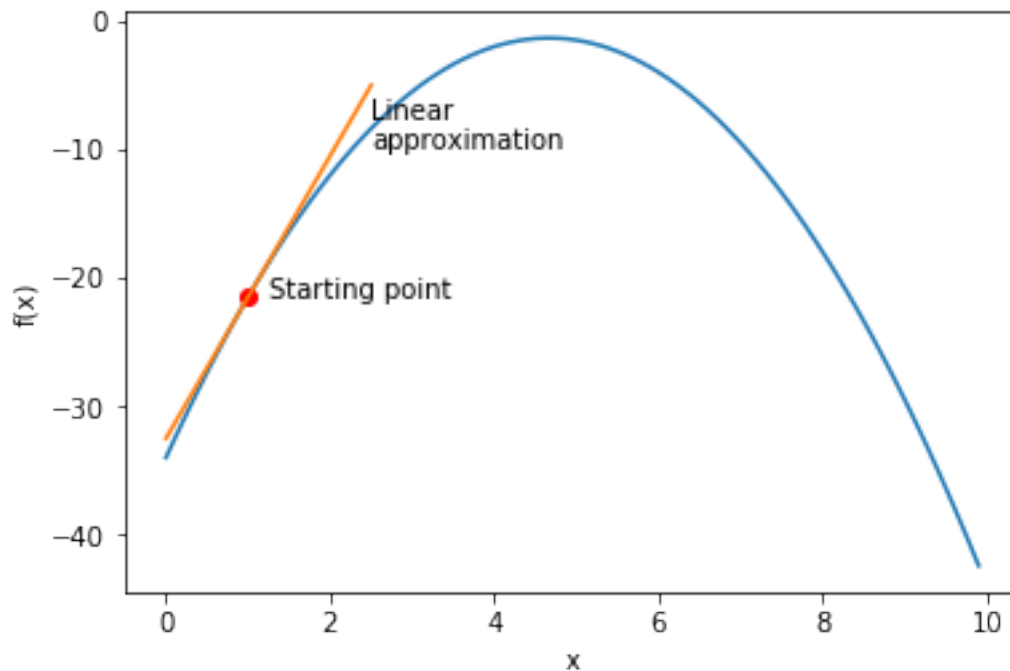
Let's first consider how we can maximize a univariate differentiable function $f(x)$ iteratively.

Q: For the function shown below, starting from the labeled point, should we increase or decrease x to get to the maximum of function $f(x)$? Why?

```
In [5]: x = numpy.arange(0.0,10.0,0.1)
        f = -(x-4.0)**2.0 - 0.5*(x-6.0)**2.0
        dfdx = -(x-4.0)*2 - (x-6.0)
        plt.plot(x,f)
        plt.xlabel('x')
        plt.ylabel('f(x)')
        plt.plot(x[10],f[10], 'ro')
        plt.annotate('Starting point', (x[10]+0.25,f[10]))
        x1 = x[10]-1
        x2 = x[10]+1.5
        f1 = f[10]-dfdx[10]*1
        f2 = f[10]+dfdx[10]*1.5
        plt.plot([x1,x2],[f1,f2])
        plt.annotate('Linear\napproximation', (x2,f2-5))
```

```
Out [5]: Text(2.5,-10,'Linear\napproximation')
```


Topographical map of Pilot Mountain



Another way to see this is to consider a very simplified version of Taylor's theorem.

Theorem. Given a function $f(\cdot)$ which is smooth at x

$$f(x + d) = f(x) + f'(x)d + O(d^2)$$

where we say that $g(x) = O(h(x))$ iff $\lim_{x \rightarrow a} \left| \frac{f(x)}{g(x)} \right| = C < +\infty$.

In words, close to x function $f(\cdot)$ is very close to being a linear function of d

$$f(x + d) = f(x) + f'(x)d$$

Slope of the best linear approximation is $f'(x)$.

$f'(x)$ tells us in which direction function grows.

Multivariate functions are a bit harder to visualize.

Hence we view them like topographical maps.

Contours connect nearby points with the same altitude (function value).


```
In [6]: import matplotlib.cm as cm
x = numpy.arange(-5.0,10.0,0.1)
y = numpy.arange(-5.0,10.0,0.1)
X,Y = numpy.meshgrid(x,y)
F = -(X-1.0)**2.0 - (Y-2.0)**2.0 + 0.5*X*Y
fig = plt.figure(figsize=plt.figaspect(0.35))
from mpl_toolkits.mplot3d import Axes3D
```

```

ax = fig.add_subplot(1, 2, 1, projection='3d')
ax.plot_surface(X,Y,F)
plt.tight_layout(6.0)
ax.set_xlabel('x')
ax.set_ylabel('y')
ax.set_zlabel('f(x,y)')
ax.plot([x[130]], [y[10]], [F[130,10]], 'ro')
ax.plot([x[70]], [y[70]], [F[70,70]], 'go')
ax.plot([x[100]], [y[30]], [F[100,30]], 'yo')
ax = fig.add_subplot(1,2,2)
ax.contour(X,Y,F,levels=numpy.arange(numpy.min(F),numpy.max(F),2))
ax.set_xlabel('x')
ax.set_ylabel('y')
ax.plot(x[130],y[10], 'ro')
ax.plot(x[70],y[70], 'go')
ax.plot(x[100],y[30], 'yo')

```

Out[6]: [<matplotlib.lines.Line2D at 0x11573e160>]



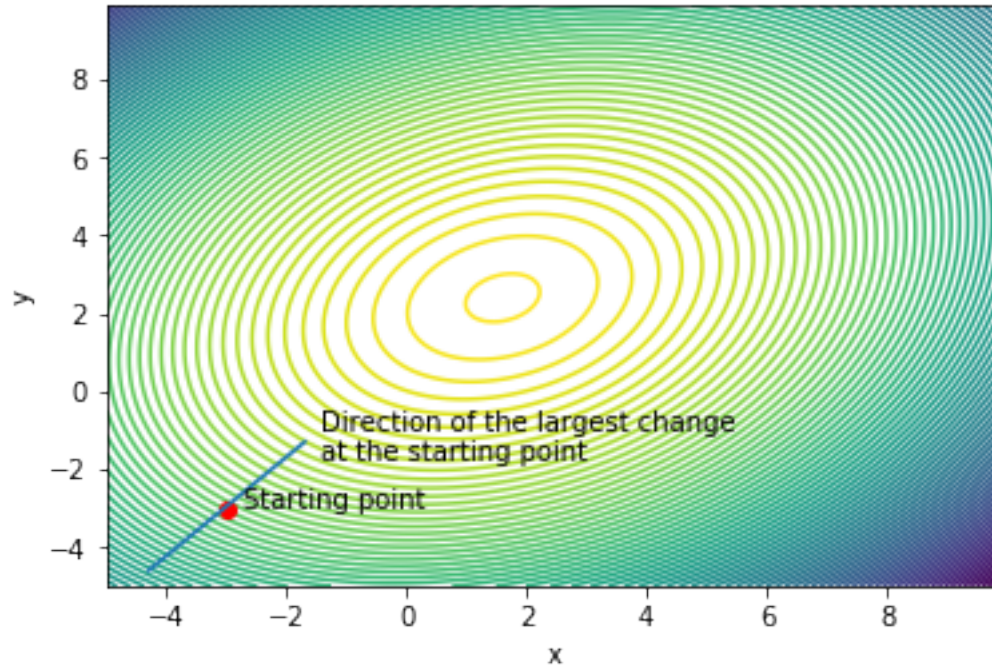
```

In [7]: import matplotlib.cm as cm
x = numpy.arange(-5.0,10.0,0.1)
y = numpy.arange(-5.0,10.0,0.1)
X,Y = numpy.meshgrid(x,y)
F = -(X-1.0)**2.0 - (Y-2.0)**2.0 + 0.5*X*Y
dfdx = -2.0*(X-1.0) + 0.5*Y
dfdy = -2.0*(Y-2.0) + 0.5*X
plt.contour(X,Y,F,levels=numpy.arange(numpy.min(F),numpy.max(F),2))
plt.annotate('Starting point', (x[20]+0.25,y[20]))
plt.plot(x[20],y[20], 'ro')
x1 = x[20] -0.2*dfdx[20,20]
y1 = y[20] -0.2*dfdy[20,10]
x2 = x[20] +0.2*dfdx[20,20]
y2 = y[20] +0.2*dfdy[20,20]
plt.annotate('Direction of the largest change \nat the starting point', (x2+0.25,y2-0.5))
plt.xlabel('x')

```

```
plt.ylabel('y')
plt.plot([x1,x2],[y1,y2])
```

Out[7]: [<matplotlib.lines.Line2D at 0x1157b4748>]



Much as in the case of univariate function, the direction in which the function changes is described by derivatives.

A vector composed of partial derivatives of a function is called gradient

$$\nabla f(\mathbf{x}) = \begin{bmatrix} \frac{\partial}{\partial x_1} f(\mathbf{x}) \\ \vdots \\ \frac{\partial}{\partial x_p} f(\mathbf{x}) \end{bmatrix}$$

For example,

$$f(x, y) = -(x - 1)^2 - (y - 2)^2 + \frac{1}{2}xy$$

has gradient

$$\nabla f(x, y) = \begin{bmatrix} -2(x - 1) + \frac{1}{2}y \\ -2(y - 2) + \frac{1}{2}x \end{bmatrix}$$

Gradient points in the direction of the largest change in the function.

Q: Can you come up with an algorithm that would use the gradients to maximize a function?

Gradients are important because: 1. They can tell us when we have reached an optimum ($\nabla f(\mathbf{x}) = 0$). 2. Point in the direction in which function is changing the most.

So, we shouldn't mess them up.

To check your gradient go back to the definition:

$$\frac{\partial}{\partial x} f(x_1, \dots, x_p) = \lim_{\delta \rightarrow 0} \frac{f(x_1 + \delta/2, \dots, x_p) - f(x_1 - \delta/2, \dots, x_p)}{\delta}$$

Left side can be numerically approximated using a small $\delta = 10^{-6}$:

$$\frac{f(x_1 + 0.5\delta, \dots, x_p) - f(x_1 - 0.5\delta, \dots, x_p)}{\delta}$$

This is called **finite difference approximation**.

Let's see this in practice.

```
In [8]: def grad_check(f,xy0,delta=1e-6,tolerance=1e-7):
    f0,g0 = f(xy0)
    p = len(xy0)
    finite_diff = numpy.zeros(p)
    gradient_correct = True
    for i in range(p):
        xy1 = numpy.copy(xy0)
        xy2 = numpy.copy(xy0)
        xy1[i] = xy1[i] - 0.5*delta
        xy2[i] = xy2[i] + 0.5*delta
        f1,_ = f(xy1)
        f2,_ = f(xy2)
        finite_diff[i] = (f2 - f1)/(delta)
        if (abs(finite_diff[i] - g0[i])>tolerance):
            print("Broken partial",i," Finite Diff: ",finite_diff[i]," Partial: ",g0[i])
            gradient_correct = False
    return gradient_correct

def f_broken(xy):
    x = xy[0]
    y = xy[1]
    f = -(x-1.0)**2.0 - (y-2.0)**2.0 + 0.5*x*y
    dfdx = -2.0*(x-1.0) + 0.5*y
    dfdy = -2.0*(y-2.0) - 0.5*x
    grad = [dfdx,dfdy]
    return f,grad

def f_correct(xy):
    x = xy[0]
    y = xy[1]
    f = -(x-1.0)**2.0 - (y-2.0)**2.0 + 0.5*x*y
    dfdx = -2.0*(x-1.0) + 0.5*y
    dfdy = -2.0*(y-2.0) + 0.5*x
    grad = [dfdx,dfdy]
    return f,grad

xy = numpy.asarray([1.0,1.0])
```

```
print("f_broken has correct gradient:", grad_check(f_broken,[1.0,1.0]))
print("f_correct has correct gradient:", grad_check(f_correct,[1.0,1.0]))
```

```
Broken partial 1 Finite Diff: 2.500000000035 Partial: 1.5
f_broken has correct gradient: False
f_correct has correct gradient: True
```

Once we are sure that the gradients are correct, we can proceed to optimize the function.

The algorithm we introduce next is called **gradient ascent** or **gradient descent** depending on whether we are maximizing or minimizing a function.

```
In [9]: def gradient_ascent(f,x,init_step,iterations):
        f_val,grad = f(x)                                # compute function value and gradient
        f_vals = [f_val]
        for it in range(iterations):                     # iterate for a fixed number of iterations
            done = False                                  # initial condition for done
            line_search_it = 0                             # how many times we tried to shrink the
            step = init_step                               # reset step size to the initial size
            while not done and line_search_it<100:         # are we done yet?
                new_x = x + step*grad                     # take a step along the gradient
                new_f_val,new_grad = f(new_x)              # evaluate function value and gradient
                if new_f_val<f_val:                        # did we go too far?
                    step = step*0.95                     # if so, shrink the step-size
                    line_search_it += 1                   # how many times did we shrink the step
                else:
                    done = True                           # better than the last x, so we move on

            if not done:                                  # did not find right step size
                print("Line Search failed.")
            else:
                f_val = new_f_val                          # ah, we are ok, accept the new x
                x = new_x
                grad = new_grad
                f_vals.append(f_val)
            plt.plot(f_vals)
        plt.xlabel('Iterations')
        plt.ylabel('Function value')
        return f_val, x
```

17 Linear Regression -- likelihood

We start by writing out a likelihood for linear regression is

$$\mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) = \prod_{i=1}^N p(y_i | \mathbf{x}_i, \beta_0, \beta, \sigma^2) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{1}{2\sigma^2} \left(y_i - (\beta_0 + \sum_j x_{ij}\beta_j) \right)^2 \right\}$$

Log-likelihood for linear regression is

$$\log \mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) = \sum_{i=1}^N \left[-\frac{1}{2} \log 2\pi\sigma^2 - \frac{1}{2\sigma^2} \left(y_i - (\beta_0 + \sum_j x_j \beta_j) \right)^2 \right].$$

18 Linear Regression -- gradient of log-likelihood

Partial derivatives

$$\begin{aligned} \frac{\partial}{\partial \beta_0} \log \mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) &= \sum_{i=1}^N -\frac{1}{\sigma^2} \left(y_i - (\beta_0 + \sum_j x_j \beta_j) \right) (-1) \\ \frac{\partial}{\partial \beta_k} \log \mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) &= \sum_{i=1}^N -\frac{1}{\sigma^2} \left(y_i - (\beta_0 + \sum_j x_j \beta_j) \right) (-x_k) \quad , k \in \{1, \dots, p\} \end{aligned}$$

We will refer to expression $y_i - (\beta_0 + \sum_j x_j \beta_j)$ as **residual**.

Hence gradient (with respect to β s)

$$\nabla \log \mathcal{L}(\beta_0, \beta, \sigma^2 | \mathbf{x}, \mathbf{y}) = \begin{bmatrix} \sum_{i=1}^N -\frac{1}{\sigma^2} \left(y_i - (\beta_0 + \sum_j x_j \beta_j) \right) (-1) \\ \sum_{i=1}^N -\frac{1}{\sigma^2} \left(y_i - (\beta_0 + \sum_j x_j \beta_j) \right) (-x_1) \\ \vdots \\ \sum_{i=1}^N -\frac{1}{\sigma^2} \left(y_i - (\beta_0 + \sum_j x_j \beta_j) \right) (-x_p) \end{bmatrix}$$

```
In [10]: X = numpy.asarray([[3.0,3.0,4.0,0.0],[3.0,3.0,3.0,0.0],[3.0,2.0,2.0,5.0],
[2.0,2.0,4.0,0.0],[3.0,3.0,4.0,4.0],[4.0,4.0,3.0,0.0],
[2.0,2.0,4.0,0.0],[2.0,2.0,2.0,0.0],[3.0,4.0,2.0,5.0],
[2.0,2.0,3.0,0.0],[2.0,4.0,2.0,0.0],[3.0,3.0,3.0,0.0],
[2.0,2.0,4.0,5.0],[4.0,3.0,2.0,6.0],[3.0,4.0,4.0,4.0],
[3.0,2.0,2.0,6.0],[3.0,3.0,3.0,0.0],[3.0,2.0,4.0,0.0],
[3.0,3.0,4.0,4.0],[3.0,4.0,2.0,0.0]])
```

```
Y = numpy.asarray([4.0,3.0,2.0,3.0,3.0,4.0,3.0,2.0,3.0,
3.0,3.0,3.0,2.0,2.0,3.0,2.0,3.0,3.0,3.0])
```

```
def linear_regression_log_likelihood(Y,X,betas,sigma2=1.0):
    ll = 0
    beta0 = betas[0]
    beta = betas[1:]
    dlldbeta0 = 0
    dlldbeta = numpy.zeros(len(beta))
    for (x,y) in zip(X,Y):
        ll = ll -0.5*numpy.log(2*numpy.pi*sigma2)
        res = y - beta0 - numpy.sum(x*beta)
        ll = ll - 1.0/(2.0*sigma2)*(res**2.0)
        dlldbeta0 = dlldbeta0 - 1.0/sigma2*res*(-1)
        dlldbeta = dlldbeta - 1.0/sigma2*(res*(-x))
```

```

grad = numpy.zeros(len(beta)+1)
grad[0] = dlldbata0
grad[1:] = dlldbata
return ll, grad

init_beta = [0.1]*5
f = lambda betas: linear_regression_log_likelihood(Y,X,betas)
grad_check(f,init_beta)

```

Out[10]: True

19 Fitting linear regression using gradient ascent

We are now ready to optimize our model.

We will aim to find maximum of the log-likelihood function using the gradient ascent algorithm we implemented.

```


In [11]: init_beta = [0.1]*5
init_step = 0.01
iterations = 1000
f = lambda betas: linear_regression_log_likelihood(Y,X,betas)
[f_best,betas_mle] = gradient_ascent(f,init_beta,0.01,100)
names = ['0','COMP410','MATH233','STOR435','Beers']
for (name,beta) in zip(names,betas_mle):
    print('Beta',name,'=',beta)

```

```

Beta 0 = 0.220839910747
Beta COMP410 = 0.278268720838
Beta MATH233 = 0.361256724745
Beta STOR435 = 0.347500790314
Beta Beers = -0.131046074777

```



20 Today

- Introduced Linear Regression
- Introduced Gradient Ascent/Descent
- Implemented Linear Regression log-likelihood and gradient ascent
- Fit a simple model for COMP 755 grades.