

# COMP755-Lect13

October 1, 2018

## 1 COMP 755

Plan for today

0. Covariance refresher
1. More Mixture Models
  - Mixture of Gaussians with Covariance
2. MapReduce
3. Debugging EM algorithms

```
In [16]: import numpy
import matplotlib.pyplot as plt
%matplotlib inline
def generate_data(N,d,K,proby,mus,As=None):
    if As is None:
        As = numpy.zeros((d,d,K))
        for k in range(K):
            As[:, :, k] = numpy.eye(d)
    ys = numpy.zeros(N, dtype='int')
    xs = numpy.zeros((d,N))
    for i in range(N):
        # Sample class according to the prior p(y)
        # in this case it is uniform
        ys[i] = numpy.random.choice(K,1)[0]
        # Sample feature values according to p(x/y)
        # In this case,  $x \sim N(\mu[y[i]], \sigma^2 I)$ 
        # To accomplish this, draw  $z_1, z_2 \sim N(0, I)$ 
        z = numpy.random.randn(2,1)
        # transform by matrix A and shift by class mean
        A = As[:, :, ys[i]].squeeze()
        mu = mus[:, ys[i]]
        Az = numpy.dot(A, z)
        x = Az + mu[:, numpy.newaxis]
        xs[:, i] = x[:, 0]
    return xs, ys
```

```

def plot_covariance(mu,Sigma,std_devs,color):
    N = 50
    alphas = numpy.linspace(0,2*numpy.pi,N)
    x = numpy.cos(alphas)
    y = numpy.sin(alphas)
    xy = numpy.vstack((x,y))
    d,v = numpy.linalg.eig(Sigma)
    d = numpy.sqrt(d)
    xy = std_devs*numpy.dot(numpy.dot(v,numpy.diag(d)),xy) + mu[:,numpy.newaxis]
    plt.plot(xy[0,:],xy[1:], 'w-',linewidth=6)
    plt.plot(xy[0,:],xy[1:],color+'-',linewidth=3)

def plot_samples(xs,ys,mus=None,Sigmas=None,colors=['r','g','b','k','c','m'],labels=None):
    N = xs.shape[1]
    if not ys is None:
        K = numpy.max(ys)+1
        for c in range(K):
            # indices of samples assigned to class c
            ind = [i for i in range(N) if ys[i]==c]
            if labels is None:
                label = "Samples in cluster " + str(c)
            else:
                label = labels[c]
            plt.plot(xs[0,ind],xs[1,ind],colors[c]+'.',label=label)
            if not mus is None:
                plt.plot(mus[0,c],mus[1,c], 'wx',markersize=9,markeredgewidth=5)
                plt.plot(mus[0,c],mus[1,c],colors[c]+'x',markersize=7,markeredgewidth=3)
            if not Sigmas is None:
                plot_covariance(mus[:,c],Sigmas[:, :,c],2.0,colors[c])
            plt.legend(loc=2, bbox_to_anchor=(1,1))
    else:
        plt.plot(xs[0,:],xs[1:],'.')

def plot_samples_post(xs,qs,mus=None,Sigmas=None, colors=['r','g','b','k','c','m'],
                    highlight_samples=None,
                    label_means=False):
    K,N = qs.shape
    for i in range(N):
        plt.plot(xs[0,i],xs[1,i], 'o',color=qs[:,i])
    if not highlight_samples is None:
        for (i,d) in highlight_samples:
            s = ''
            for j in range(3):
                if j>0:
                    s = s + '\n'
                s = s + '$p(h_{\{\{\}\}}=\{\} |x_{\{\{\}\}})=\${:1.4f})'.format(i,j,i,qs[j,i])
            if d==0:

```

```

        plt.annotate(s,xy=(xs[0,i]+0.5,xs[1,i]-1.0),
                     bbox=dict(facecolor='white'),
                     fontsize=15)
        plt.arrow(xs[0,i]+0.5,xs[1,i],-0.5,0)
    if d==1:
        plt.annotate(s,xy=(xs[0,i]-2.0,xs[1,i]-4.5),
                     bbox=dict(facecolor='white'),
                     fontsize=15)
        plt.arrow(xs[0,i],xs[1,i]-4.5,0,4.5)
for c in range(K):
    # indices of samples assigned to class c
    if not mus is None:
        plt.plot(mus[0,c],mus[1,c], 'kx',markersize=9,markeredgewidth=5)
        plt.plot(mus[0,c],mus[1,c], 'wx',markersize=7,markeredgewidth=3)
    if label_means:
        plt.annotate('$\mu_{\text{' + str(c) + '}}$',xy=(mus[0,c]+1,mus[1,c]-1),
                     bbox=dict(facecolor='white'),
                     fontsize=15)
    if not Sigmas is None:
        plot_covariance(mus[:,c],Sigmas[:, :, c],2.0,colors[c])

```

In [17]: `import matplotlib.pyplot as plt`

```

plt.figure(figsize=(10,10))
plt.subplot(2,2,1)
K = 3
d=2
mus = 10*numpy.asarray([[0.0,1.0,2.0],[0.0,0.0,0.0]])
As = numpy.asarray([[2.0,0.0],[0.0,2.0]],
                    [[0.5,0.0],[0.0,5.0]],
                    [[2.0,-2.0],[0.0,2.0]])
As = numpy.swapaxes(As,0,2)
Sigmas = numpy.zeros((d,d,K))
for c in range(K):
    A = As[:, :, c]
    Sigmas[:, :, c] = numpy.dot(A,A.transpose())
    print "Sigma"+str(c)
    print Sigmas[:, :, c]

proby = [1./K]*K
numpy.random.seed(1)
xs,ys = generate_data(1000,2,K,proby,mus,As)

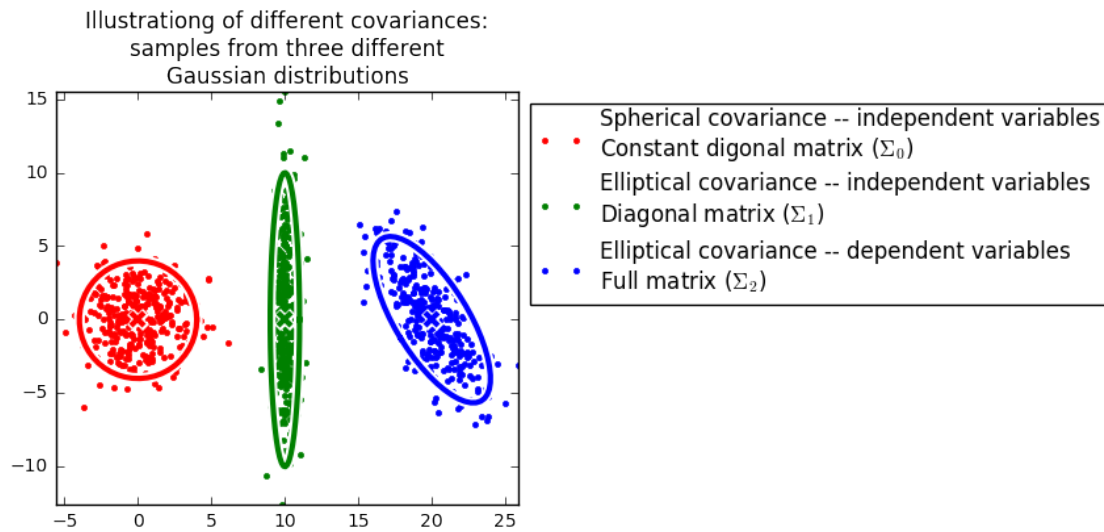
plot_samples(xs,ys,mus=mus,Sigmas=Sigmas,
             labels=['Spherical covariance -- independent variables\nConstant digonal m
                    'Elliptical covariance -- independent variables\nDiagonal matrix
                    'Elliptical covariance -- dependent variables\nFull matrix ($\text{Si}

plt.axis('image')

```

```
plt.title('Illustration of different covariances:\nsamples from three different\nGauss')

Sigma0
[[ 4.  0.]
 [ 0.  4.]]
Sigma1
[[ 0.25  0. ]
 [ 0.    25. ]]
Sigma2
[[ 4. -4.]
 [-4.  8.]]
```



## 2 Our second EM algorithm

The model

$$p(h \mid \alpha) = \alpha_h$$

$$p(\mathbf{x} \mid h, \mu) = (2\pi)^{-\frac{d}{2}} |\Sigma_h|^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \mu_{h_t})^T \Sigma_h^{-1} (\mathbf{x} - \mu_{h_t}) \right\}$$

is a variant of \*\* Mixture of Gaussians. \*\* Note that we introduced a covariance matrix per cluster.

- Hidden variables:  $h_t$  -- cluster membership for sample  $t$
- Parameters:  $\Theta = (\underbrace{\alpha_1, \dots, \alpha_K}_{\text{proportions}}, \underbrace{\mu_1, \dots, \mu_K}_{\text{means}}, \underbrace{\Sigma_1, \dots, \Sigma_K}_{\text{covariances}})$

### 3 Our second EM algorithm

We plug-in probabilities  $p(\mathbf{x}_t \mid h_t, \Theta)$  and  $p(h_t \mid \alpha)$  in the bound

$$\begin{aligned} \mathcal{B}(\Theta, q) &= \sum_{t=1}^T \sum_{h_t} q_t(h_t) \log \frac{p(\mathbf{x}_t, h_t \mid \Theta)}{q_t(h_t)} \\ &= \sum_{t=1}^T \sum_{h_t} q_t(h_t) \left[ \log \alpha_{h_t} - \frac{d}{2} \log(2\pi) - \frac{1}{2} \log |\Sigma_{h_t}| \right. \\ &\quad \left. - \frac{1}{2} (\mathbf{x}_t - \mu_{h_t})^T \Sigma_{h_t}^{-1} (\mathbf{x}_t - \mu_{h_t}) \right] \\ &\quad - \sum_{t=1}^T \sum_{h_t} q_t(h_t) \log q_t(h_t) \end{aligned}$$

```
In [18]: # broadcasting tutorial
x = numpy.asmatrix([[1,2,3],[4,5,6]])
print "Data matrix:"
print x
mu = numpy.asarray([1,2])
print "Mean as a row vector:"
print mu
print "Mean as a column vector:"
print mu[:,numpy.newaxis]
print "Broadcast subtraction across columns:"
print x - mu[:,numpy.newaxis]
```

Data matrix:

```
[[1 2 3]
 [4 5 6]]
```

Mean as a row vector:

```
[1 2]
```

Mean as a column vector:

```
[[1]
 [2]]
```

Broadcast subtraction across columns:

```
[[0 1 2]
 [2 3 4]]
```

## 4 Our second EM algorithm -- E-step

The E-step

$$\begin{aligned}
 q_t(h_t = k) &= p(h_t = k \mid \mathbf{x}_t, \mu) = \frac{p(\mathbf{x}_t, h_t = k \mid \mu)}{\underbrace{\sum_c p(\mathbf{x}_t, h_t = c \mid \mu)}_{\text{same for all values of } k}} \\
 &\propto p(\mathbf{x}_t, h_t = k \mid \mu) \\
 &= \alpha_{h_t} (2\pi)^{-\frac{d}{2}} |\Sigma_{h_t}^{-1}| \exp \left\{ -\frac{1}{2} (\mathbf{x} - \mu_{h_t})^T \Sigma_{h_t}^{-1} (\mathbf{x} - \mu_{h_t}) \right\}
 \end{aligned}$$

Implementation:

```

q = numpy.zeros((K,N))          # clusters x samples
q = logjointp(x,Theta)          # compute all joints at once
loglik = numpy.sum(logsumexp(q)) # compute loglikelihood
q = q - logsumexp(q)            # normalizing across clusters

In [19]: def logjointp(xs,mus,Sigmas,logph = None):
    # compute log p(x,h/mus,Sigmas)
    # for all values of h (1..K)
    d = xs.shape[0]
    N = xs.shape[1]
    K = mus.shape[1]
    # dimensions are consistent
    assert(mus.shape[0] == d and
           Sigmas.shape[0] == d and
           Sigmas.shape[1] == d and
           Sigmas.shape[2] == K)
    if logph is None:
        # no prior for clusters provided
        # assume uniform probability
        logph = numpy.asarray([1./K]*K)

    logp = numpy.zeros( (K,N) )
    for k in range(K):
        # get the covaraiance and squeeze
        # the last dimension to get d x d matrix
        Sigma = Sigmas[:, :, k].squeeze()
        mu = mus[:, k]
        invSigma = numpy.linalg.inv(Sigma)
        ###
        # For each sample we need to compute:
        # (x - mu) Sigma^{-1} (x - mu)
        # We will do this at once for all samples
        #
        # (x-mu)^T for all rows of matrix xs
        # mu[:,numpy.newaxis] turns row vector into a column
        res = xs - mu[:,numpy.newaxis]

```

```

# (x-mu)^T Sigma^{-1} for all rows of matrix X
res_T_invSigma = numpy.dot(res.transpose(),invSigma)
# need to compute inner product between
# *corresponding* rows of ((x-mu)^T Sigma^{-1}) and (x-mu)
res_T_invSigma_res = numpy.sum(res_T_invSigma.transpose()*res,axis=0)
logp[k,:] = -0.5*res_T_invSigma_res
logp[k,:] = logp[k,:] - d/2.*numpy.log(2.*numpy.pi)
logp[k,:] = logp[k,:] - 0.5*numpy.log(numpy.linalg.det(Sigma))
logp[k,:] = logp[k,:] + logph[k]
return logp

```

## 5 Our second EM algorithm -- M-step

Updates for parameters of prior probability  $p(h | \alpha)$

$$\alpha_c^* = \frac{\sum_t q_t(h_t = c)}{N}$$

Updates for means of clusters

$$\mu_c^* = \frac{\sum_t q_t(h_t = c) \mathbf{x}_t}{\sum_t q_t(c)}$$

Updates for covariances of clusters

$$\Sigma_c^* = \frac{\sum_t q_t(h_t = c) (\mathbf{x}_t - \mu_c^*)(\mathbf{x}_t - \mu_c^*)^T}{\sum_t q_t(c)}$$

Work out means and covariances on the board.

## 6 Matrix calculus

$$\begin{aligned}
\nabla_{\mathbf{a}} \mathbf{b}^T \mathbf{a} &= \mathbf{b} \\
\nabla_{\mathbf{a}} \mathbf{a}^T \mathbf{A} \mathbf{a} &= (\mathbf{A} + \mathbf{A}^T) \mathbf{a} \\
\nabla_{\mathbf{A}} \text{tr} \{ \mathbf{B} \mathbf{A} \} &= \mathbf{B}^T \\
\nabla_{\mathbf{A}} \log |\mathbf{A}| &= (\mathbf{A}^{-1})^T \\
\nabla_{\mathbf{A}} -\log |\mathbf{A}| &= (\mathbf{A})^T \\
\text{tr} \{ \mathbf{A} + \mathbf{B} \} &= \text{tr} \{ \mathbf{A} \} + \text{tr} \{ \mathbf{B} \} \\
\text{tr} \{ \mathbf{A} \mathbf{B} \mathbf{C} \} &= \text{tr} \{ \mathbf{C} \mathbf{A} \mathbf{B} \} = \text{tr} \{ \mathbf{B} \mathbf{C} \mathbf{A} \} \\
\mathbf{x}^T \mathbf{A} \mathbf{x} &= \text{tr} \{ \mathbf{x}^T \mathbf{A} \mathbf{x} \} = \text{tr} \{ \mathbf{x} \mathbf{x}^T \mathbf{A} \} = \text{tr} \{ \mathbf{A} \mathbf{x} \mathbf{x}^T \}
\end{aligned}$$

where  $\mathbf{a}$  and  $\mathbf{b}$  are vectors;  $\mathbf{A}$ ,  $\mathbf{B}$ , and  $\mathbf{C}$  are matrices;  $\text{tr} \{ \mathbf{A} \} = \sum_i \mathbf{A}_{i,i}$  is trace of matrix.

We will use the fact that covariance matrices and their inverses are symmetric  $\Sigma = \Sigma^T, \Sigma^{-1} = \Sigma^{-T}$

Refer to section 4.1.3. of your textbook.

## 7 Our second EM algorithm -- details of M-step derivation

Bound:

$$\begin{aligned}
\mathcal{B}(\Theta, q) &= \sum_{t=1}^T \sum_{h_t} q_t(h_t) \log \frac{p(\mathbf{x}_t, h_t \mid \Theta)}{q_t(h_t)} \\
&= \sum_{t=1}^T \sum_{h_t} q_t(h_t) \left[ \log \alpha_{h_t} - \frac{d}{2} \log(2\pi) - \frac{1}{2} \log |\Sigma_{h_t}| \right. \\
&\quad \left. - \frac{1}{2} (\mathbf{x}_t - \mu_{h_t})^T \Sigma_{h_t}^{-1} (\mathbf{x}_t - \mu_{h_t}) \right] \\
&\quad - \sum_{t=1}^T \sum_{h_t} q_t(h_t) \log q_t(h_t)
\end{aligned}$$

## 8 Our second EM algorithm -- details of M-step derivation

Parts of bound relevant for updates of  $\mu_c$

$$\sum_{t=1}^T q_t(h_t = c) \left[ -\frac{1}{2} (\mathbf{x}_t - \mu_c)^T \Sigma_c^{-1} (\mathbf{x}_t - \mu_c) \right]$$

Parts of bound relevant for updates of  $\Sigma_c$

$$\sum_{t=1}^T q_t(h_t = c) \left[ -\frac{1}{2} \log |\Sigma_c| - \frac{1}{2} (\mathbf{x}_t - \mu_c)^T \Sigma_c^{-1} (\mathbf{x}_t - \mu_c) \right]$$

We replace  $\Sigma_c^{-1}$  with  $\Lambda_c$ . Parts of bound relevant for updates of  $\Lambda_c$

$$\sum_{t=1}^T q_t(h_t = c) \left[ \frac{1}{2} \log |\Lambda_c| - \frac{1}{2} (\mathbf{x}_t - \mu_c)^T \Lambda_c (\mathbf{x}_t - \mu_c) \right]$$

## 9 Matrix calculus

$$\begin{aligned}
\nabla_{\mathbf{a}} \mathbf{b}^T \mathbf{a} &= \mathbf{b} \\
\nabla_{\mathbf{a}} \mathbf{a}^T \mathbf{A} \mathbf{a} &= (\mathbf{A} + \mathbf{A}^T) \mathbf{a} \\
\nabla_{\mathbf{A}} \text{tr} \{ \mathbf{B} \mathbf{A} \} &= \mathbf{B}^T \\
\nabla_{\mathbf{A}} \log |\mathbf{A}| &= (\mathbf{A}^{-1})^T \\
\nabla_{\mathbf{A}} -\log |\mathbf{A}| &= (\mathbf{A})^T \\
\text{tr} \{ \mathbf{A} + \mathbf{B} \} &= \text{tr} \{ \mathbf{A} \} + \text{tr} \{ \mathbf{B} \} \\
\text{tr} \{ \mathbf{A} \mathbf{B} \mathbf{C} \} &= \text{tr} \{ \mathbf{C} \mathbf{A} \mathbf{B} \} = \text{tr} \{ \mathbf{B} \mathbf{C} \mathbf{A} \} \\
\mathbf{x}^T \mathbf{A} \mathbf{x} &= \text{tr} \{ \mathbf{x}^T \mathbf{A} \mathbf{x} \} = \text{tr} \{ \mathbf{x} \mathbf{x}^T \mathbf{A} \} = \text{tr} \{ \mathbf{A} \mathbf{x} \mathbf{x}^T \}
\end{aligned}$$

where  $\mathbf{a}$  and  $\mathbf{b}$  are vectors;  $\mathbf{A}$ ,  $\mathbf{B}$ , and  $\mathbf{C}$  are matrices;  $\text{tr} \{ \mathbf{A} \} = \sum_i \mathbf{A}_{i,i}$  is trace of matrix.

We will use the fact that covariance matrices and their inverses are symmetric  $\Sigma = \Sigma^T, \Sigma^{-1} = \Sigma^{-T}$

Refer to section 4.1.3. of your textbook.



## 10 Our second EM algorithm -- details of $\mu_c$ update

Equate gradient to 0

$$\nabla_{\mu_c} \sum_{t=1}^T q_t(h_t = c) \left[ -\frac{1}{2} (\mathbf{x}_t - \mu_c)^T \Lambda (\mathbf{x}_t - \mu_c) \right] = 0$$

Expand the quadratic term and move  $-\frac{1}{2}$

$$\nabla_{\mu_c} \sum_{t=1}^T q_t(h_t = c) \left( -\frac{1}{2} \right) \left[ \mathbf{x}_t^T \Lambda_c \mathbf{x}_t - \mathbf{x}_t^T \Lambda_c \mu_c - \mu_c^T \Lambda_c \mathbf{x}_t + \mu_c^T \Lambda_c \mu_c \right] = 0$$

gradient of sum is sum of gradients

$$\sum_{t=1}^T q_t(h_t = c) \left( -\frac{1}{2} \right) \underbrace{\nabla_{\mu_c} \left[ \mathbf{x}_t^T \Lambda_c \mathbf{x}_t - \mathbf{x}_t^T \Lambda_c \mu_c - \mu_c^T \Lambda_c \mathbf{x}_t + \mu_c^T \Lambda_c \mu_c \right]}_{\text{per sample contribution to the gradient}} = 0$$

## 11 Our second EM algorithm -- details of $\mu_c$ update

Focus on a single sample's contribution to the gradient

$$\begin{aligned} \nabla_{\mu_c} & \left[ \mathbf{x}_t^T \Lambda_c \mathbf{x}_t - \mathbf{x}_t^T \Lambda_c \mu_c - \mu_c^T \Lambda_c \mathbf{x}_t + \mu_c^T \Lambda_c \mu_c \right] \\ &= \left[ 0 - \Lambda_c^T \mathbf{x}_t - \mathbf{x}_t^T \Lambda_c + (\Lambda_c + \Lambda_c^T) \mu_c \right] \\ &= \left[ 0 - \Lambda_c \mathbf{x}_t - \mathbf{x}_t^T \Lambda_c + (\Lambda_c + \Lambda_c) \mu_c \right] \\ &= \left[ 2\Lambda_c^T (\mathbf{x}_t - \mu_c) \right] \end{aligned}$$

plug this back into the gradient

$$\begin{aligned} \nabla_{\mu_c} \mathcal{B}(\theta, q) &= \sum_{t=1}^T q_t(h_t = c) \left( -\frac{1}{2} \right) \left[ 2\Lambda_c^T (\mathbf{x}_t - \mu_c) \right] \\ &= -\Lambda_c^T \sum_{t=1}^T q_t(h_t = c) (\mathbf{x}_t - \mu_c) \end{aligned}$$

Equating the gradient to zero yields

$$\mu_c^* = \frac{\sum_{t=1}^T q_t(h_t = c) \mathbf{x}_t}{\sum_{t=1}^T q_t(h_t = c)}$$

## 12 Our second EM algorithm -- gory details of $\Sigma_c$ update

Start with parts of the bound relevant to  $\Lambda_c$

$$\sum_{t=1}^T q_t(h_t = c) \left[ \frac{1}{2} \log |\Lambda_c| - \frac{1}{2} (\mathbf{x}_t - \mu_c)^T \Lambda_c (\mathbf{x}_t - \mu_c) \right]$$

First use the trace permutation trick  $\mathbf{x}^T \mathbf{A} = \text{tr} \{ \mathbf{x} \mathbf{x}^T \mathbf{A} \}$

$$\sum_{t=1}^T q_t(h_t = c) \left[ \frac{1}{2} \log |\Lambda_c| - \frac{1}{2} \text{tr} \left\{ (\mathbf{x}_t - \mu_c)(\mathbf{x}_t - \mu_c)^T \Lambda_c \right\} \right]$$

Distribute sum and

$$\frac{1}{2} \log |\Lambda_c| \underbrace{\sum_{t=1}^T q_t(h_t = c)}_{w_c} - \sum_{t=1}^T q_t(h_t = c) \frac{1}{2} \text{tr} \left\{ \Lambda_c (\mathbf{x}_t - \mu_c)(\mathbf{x}_t - \mu_c)^T \right\}$$

use linearity of trace

$$\frac{1}{2} \log |\Lambda_c| w_c - \frac{1}{2} \text{tr} \left\{ \Lambda_c \underbrace{\left[ \sum_{t=1}^T q_t(h_t = c) (\mathbf{x}_t - \mu_c)(\mathbf{x}_t - \mu_c)^T \right]}_{\mathbf{S}_c} \right\}$$

to obtain

$$\frac{w_c}{2} \log |\Lambda_c| - \frac{1}{2} \text{tr} \{ \Lambda_c \mathbf{S}_c \}$$

### 13 Our second EM algorithm -- less gory details of $\Sigma_c$ update

Terms in bound relevant to update of  $\Lambda_c$  and consequently  $\Sigma_c$

$$\frac{1}{2} \log |\Lambda_c| w_c - \frac{1}{2} \text{tr} \{ \Lambda_c \mathbf{S}_c \}$$

where

$$w_c = \sum_{t=1}^T q_t(h_t = c)$$

$$\mathbf{S}_c = \sum_{t=1}^T q_t(h_t = c) (\mathbf{x}_t - \mu_c)(\mathbf{x}_t - \mu_c)^T$$

Gradient of bound

$$\begin{aligned} \nabla_{\Lambda_c} \mathcal{B}(q, \Theta) &= \nabla_{\Lambda_c} \left[ \frac{w_c}{2} \log |\Lambda_c| - \frac{1}{2} \text{tr} \{ \Lambda_c \mathbf{S}_c \} \right] \\ &= \frac{w_c}{2} \Lambda_c^{-T} - \frac{1}{2} \mathbf{S}_c^T \end{aligned}$$

Since  $\Lambda_c^{-1} = \Sigma_c$

$$\Sigma_c = \frac{\mathbf{S}_c}{w_c} = \frac{\sum_{t=1}^T q_t(h_t = c) (\mathbf{x}_t - \mu_c)(\mathbf{x}_t - \mu_c)^T}{\sum_{t=1}^T q_t(h_t = c)}$$

### 14 So much math ...

Updates for parameters of prior probability  $p(h | \alpha)$

$$\alpha_c^* = \frac{\sum_t q_t(h_t = c)}{N}$$

Updates for means of clusters

$$\mu_c^* = \frac{\sum_t q_t(c) \mathbf{x}_t}{\sum_t q_t(c)}$$

Updates for covariances of clusters

$$\Sigma_c^* = \frac{\sum_t q_t(c) (\mathbf{x}_t - \mu_c^*)(\mathbf{x}_t - \mu_c^*)^T}{\sum_t q_t(c)}$$

## 15 Updates for axis aligned covariances

A full covariance matrix has quadratically many parameters.

In order to simplify things we can assume diagonal covariance matrix and just learn per feature variance -- diagonal covariance matrix

Model:

$$p(h \mid \alpha) = \alpha_h$$

$$p(\mathbf{x} \mid h, \mu, \sigma) = (2\pi)^{-\frac{d}{2}} \frac{1}{\prod_i \sigma_{i,h}} \exp \left\{ - \sum_{i=1}^d \frac{1}{2\sigma_{i,h}^2} (x_i - \mu_{i,h})^2 \right\}$$

Updates for  $\sigma_{i,c}$

$$\sigma_{c,i}^* = \frac{\sum_t q_t(c) (x_{i,t} - \mu_{i,c}^*)^2}{\sum_t q_t(c)}$$

```
In [5]: def logsumexp(vec):
        m = numpy.max(vec,axis=0)
        return numpy.log(numpy.sum(numpy.exp(vec-m),axis=0))+m

def mog(xs,K,iterations=10, visualize=False):
    d,N = xs.shape
    # compute mean and std of data
    data_mean = numpy.mean(xs,axis=1)
    data_std = numpy.std(xs,axis=1)
    # initialize means around the data mean but
    # ensure they are not exactly the same by adding
    # small amount of noise
    mus = (data_mean[:,numpy.newaxis] +
            0.1*data_std[:,numpy.newaxis]*numpy.random.randn(d,K))
    Sigmas = numpy.zeros((d,d,K))
    for k in range(K):
        # start with large covariance
        Sigmas[:, :, k] = 0.1*numpy.eye(d) + 10.*numpy.diag(data_std)

    # assume uniform prior
    logph = numpy.array([-numpy.log(K)]*K)

    logliks = []
    for it in range(iterations):
```

```

# E-step
q = logjointp(xs,mus,Sigmas,logph)
assert(q.shape[0] == K)
loglik = numpy.sum(logsumexp(q))
logliks.append(loglik)
q = numpy.exp(q - logsumexp(q))

# M-step:
mus = numpy.dot(xs,q.transpose()/(1e-5 + numpy.sum(q,axis=1)))
logph = numpy.log(numpy.sum(q,axis=1)/N)
for k in range(K):
    mu = mus[:,k]
    res = xs - mu[:,numpy.newaxis]
    Sigma = numpy.dot(q[k,:]*res,res.transpose())/numpy.sum(q[k,:])
    Sigmas[:, :,k] = Sigma

if visualize and it % 10 == 0:
    print "Iteration: {} Log-likelihood: {}".format(it,loglik)
    plt.figure()

    plot_samples_post(xs,q,mus,Sigmas,label_means=it>0)
    plt.title(('Iteration {} Log-likelihood {} \n '+'
              'Red intensity = $p(h=0|x)$\n'+
              'Green intensity = $p(h=1|x)$\n' +
              'Blue intensity = $p(h=2|x)$').format(it,loglik),
              multialignment='right')

plt.figure()
plt.plot(logliks)
plt.xlabel('Iterations')
plt.ylabel('Log-likelihood')
alphas = numpy.exp(logph)
return mus,alphas,q

```

```

mog(xs,3,iterations=50,visualize=True);

```

```

Iteration: 0 Log-likelihood: -6388.48031418
Iteration: 10 Log-likelihood: -6154.91050602
Iteration: 20 Log-likelihood: -6092.84850651
Iteration: 30 Log-likelihood: -5784.20738195
Iteration: 40 Log-likelihood: -5112.58748288

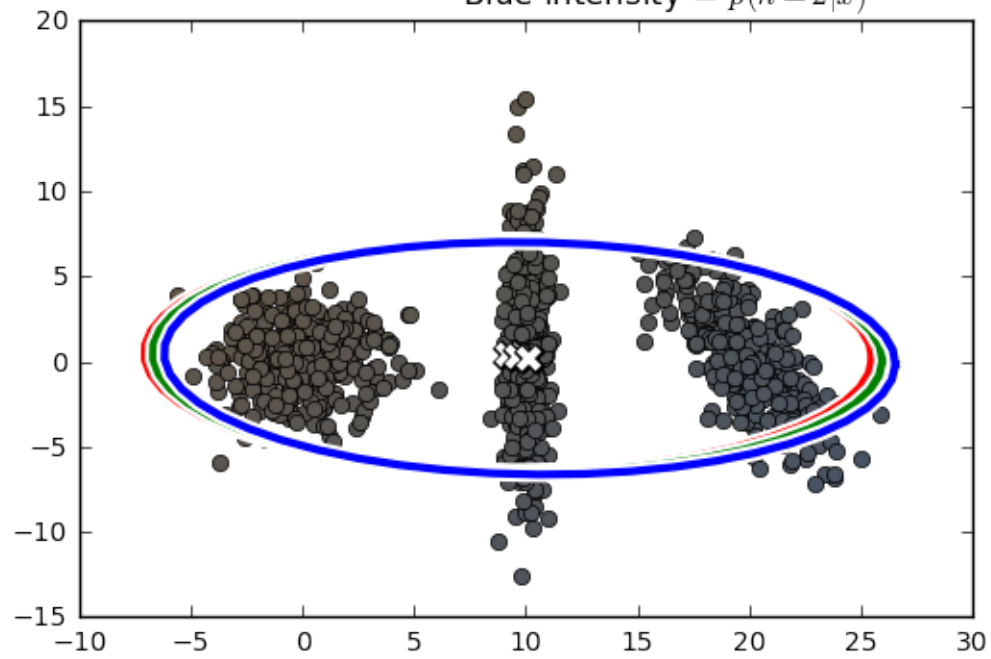
```

Iteration 0 Log-likelihood -6388.48031418

Red intensity =  $p(h = 0|x)$

Green intensity =  $p(h = 1|x)$

Blue intensity =  $p(h = 2|x)$

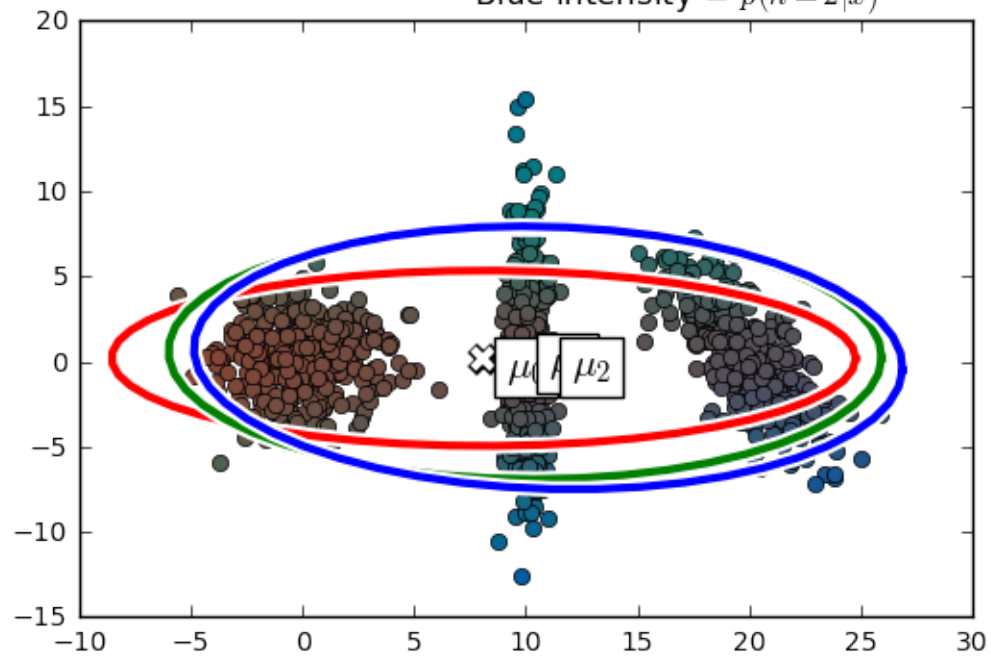


Iteration 10 Log-likelihood -6154.91050602

Red intensity =  $p(h=0|x)$

Green intensity =  $p(h=1|x)$

Blue intensity =  $p(h=2|x)$

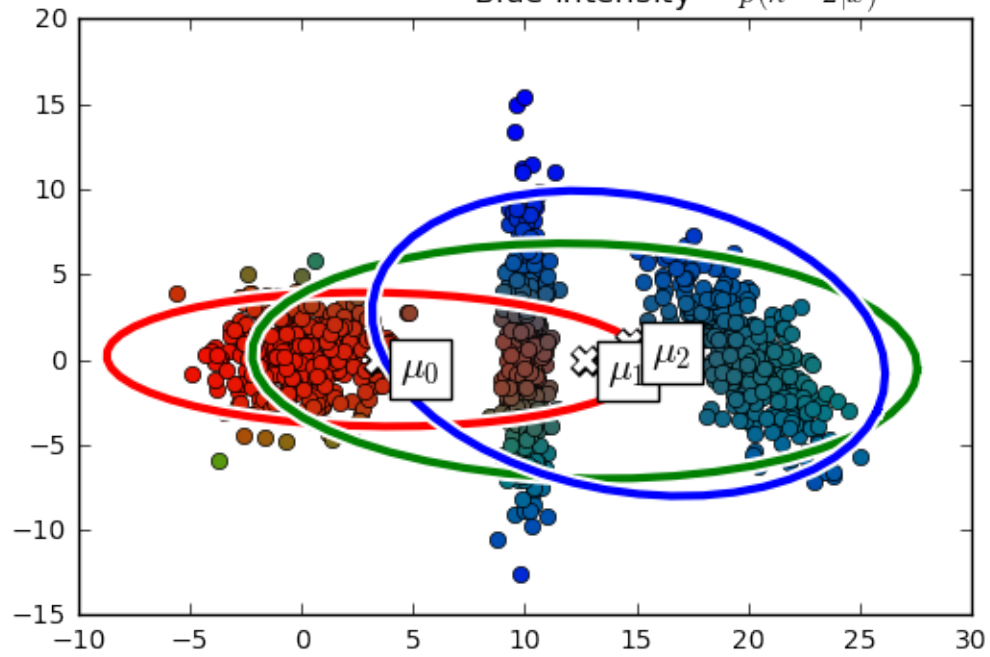


Iteration 20 Log-likelihood -6092.84850651

Red intensity =  $p(h=0|x)$

Green intensity =  $p(h=1|x)$

Blue intensity =  $p(h=2|x)$

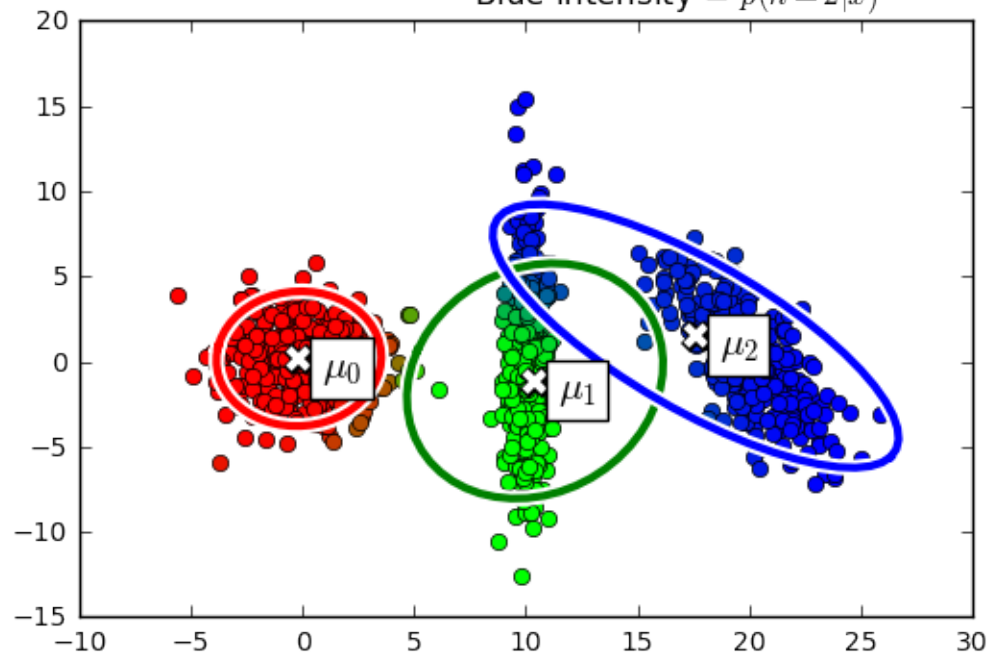


Iteration 30 Log-likelihood -5784.20738195

Red intensity =  $p(h=0|x)$

Green intensity =  $p(h=1|x)$

Blue intensity =  $p(h=2|x)$



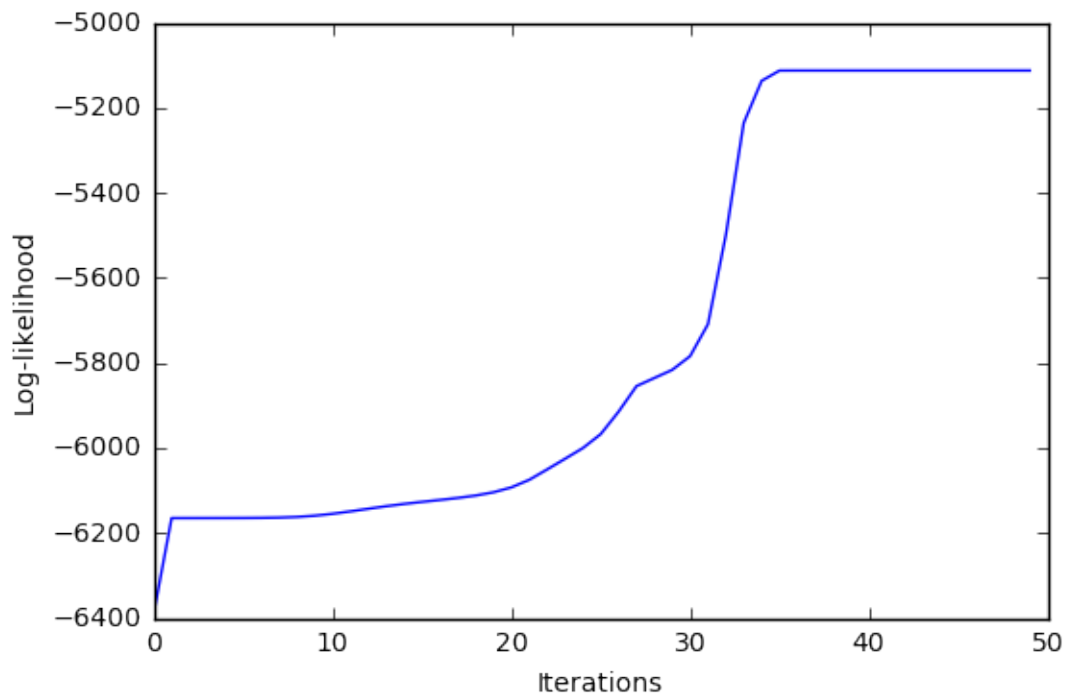
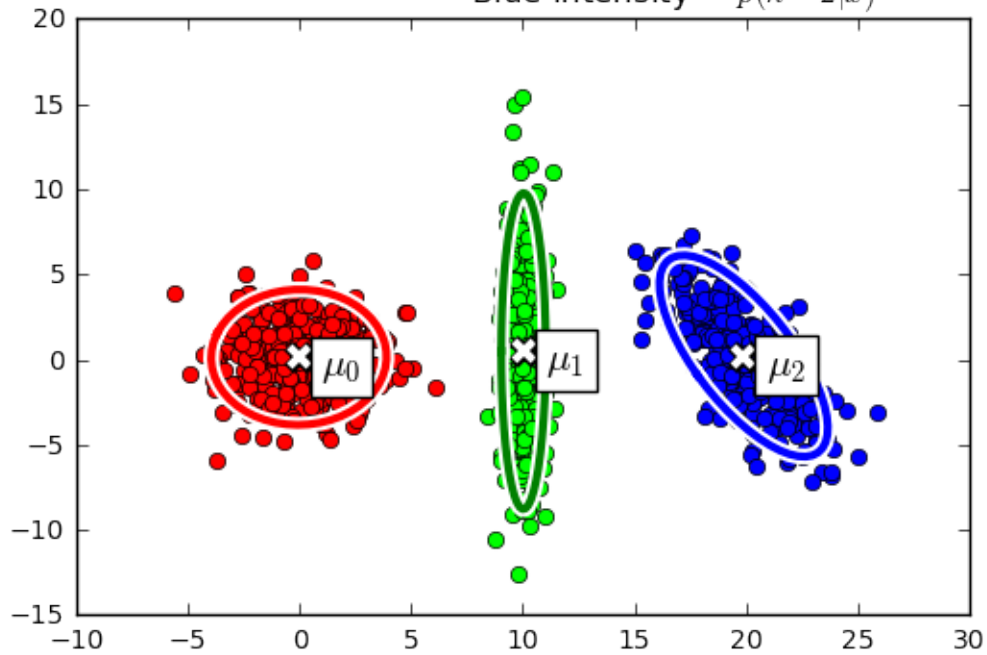


Iteration 40 Log-likelihood -5112.58748288

Red intensity =  $p(h=0|x)$

Green intensity =  $p(h=1|x)$

Blue intensity =  $p(h=2|x)$



## 16 Debugging EM

1. Log-likelihood should always go up!
2. Synthetic data is your friend. If you generate data from your model you get samples and cluster membership.
3. E-step computes cluster membership based on parameters. Use this!
  - Synthesize data from ground truth parameters
  - Start your EM from ground truth parameters, not random initialization
  - Does your E step associate samples with correct clusters?
  - Select one sample and look at its posterior probability for the cluster it came from
4. M-step updates parameters based on cluster membership. Use this!
  - Using synthetic data, set  $q$  to be one-hot according to ground truth
  - Start your M-step with this  $q$
  - If you don't get parameters back that are close to the ground truth
  - To isolate a broken update, let M-step update just one parameter (for example  $\mu$ )
5. Starting your EM with ground truth parameters should not budge too much.

Between these tricks you should be able to isolate source of your problem.

## 17 MapReduce

`map` is a common function in functional programming languages

Here is an implementation in python

```
def map(f,lst):  
    return [f(v) for v in lst]
```

`fold` or `reduce` is its companion

```
def reduce(f,lst,a):  
    if len(lst) == 0  
        return a  
    return f(lst[0],reduce(f,lst[1:],a))
```

```
In [15]: # simple map and reduce examples  
def map(f,lst):  
    return [f(v) for v in lst]  
  
def reduce(f,lst,a):  
    if len(lst) == 0:  
        return a  
    return f(lst[0],reduce(f,lst[1:],a))
```

```

def square(x):
    return x**2

def add(x,y):
    return x+y
def square(x):
    return x**2

lst = [1,2,3,4,5]
sqlist = map(square, [1,2,3,4,5])
print "Sum of squares map-reduce style:", reduce(add,sqlist,0)
print "Sum of squares numpy style:", numpy.sum(numpy.asarray(lst)**2)

```

Sum of squares map-reduce style: 55  
Sum of squares numpy style: 55

## 18 MapReduce for EM

1. map applies a function to each entry in a list
  - in our example: squaring
2. reduce summarizes the resulting list
  - in our example: sum

In the case of EM algorithm 1. In E-step, for each sample we compute  $q(h_t) = p(h_t | \mathbf{x}_t, \Theta)$  --  
map 2. In M-step, we aggregate data  $\mu_c^* = \frac{\sum_i q_i(c) \mathbf{x}_i}{\sum_i q_i(c)}$  -- reduce

The main point here is that both map and reduce phase can be divided into subtasks \* compute  $q(h)$  for subsets of data \* aggregate weighted sums for subsets of data

## 19 MapReduce for EM

Hence, EM permits trivial parallelization.

MapReduce, despite the fact that it is a standard func. programming concept and implemented in various guises all over the place, is patented by Google.

Regardless, it is a good idea to take note of the parallelization opportunities.

Hadoop is a popular and robust open source implementation.

## 20 Covered

- Mixture of Gaussian with Covariances
- Details of update derivation for means, covariances
- Debugging EM
- MapReduce